



KONG AND AWS PARTNERCAST

Build, Secure & Scale Agentic AI with Kong AI Gateway and Amazon Bedrock

Robin Cher

APJ Solutions Architect
At Kong

robin.cher@konghq.com

William Garcia

APJ Infra Solutions Architect
At Amazon Web Services

willgrod@amazon.com



AWS ISV Workload Migration, ISV Accelerate

AWS SaaS Factory, Global Startup

Amazon EKS / Graviton / Linux / Lambda Ready

Kong recognizes **AWS as Winner for the Global Partner of the Year award in 2025**

Kong serves as a critical enabler and accelerator for enterprises in their journey to **embrace AI, modernize legacy workloads and successfully scale production on AWS.**

Compliance and data transfer cost savings with private connectivity via **AWS PrivateLink**

Flexible deployment options on AWS and 20+ Native Integrations with **AWS Containers and Kong plugins for AI-based APIs**

**Kong Accelerates Migration, Modernization,
Generative AI and Agentic AI on AWS**

Trusted by Organizations of All Types

Kong powers APIs from all over the world

1000+

ENTREPRISE CUSTOMERS

28

COUNTRIES

#1

API GATEWAY

CONGLOMERATES



SIEMENS

Honeywell

Panasonic



SOFTWARE & TECHNOLOGY



S&P

ELSEVIER



NetApp

TELECOMM & ENTERTAINMENT



true



FINANCIAL SERVICES

Goldman Sachs

JPMORGAN CHASE & CO.



HSBC



Santander



DBS

HEALTHCARE

NOVARTIS



moderna



ThermoFisher SCIENTIFIC



RETAIL



NORDSTROM



CHEILJEDANG





End-to-end service connectivity

Service decentralization and collaboration within apps, between apps, and at the edge



Architectural flexibility

Connection with AWS-backed backends eg – AWS DB, Event/Streaming and AI services



Multi-environment & multi-region

APIs deployed across on-premise and AWS in migration, regulatory and disaster recovery contexts



Kubernetes native

Single operating environment for containers, microservices, and APIs on Amazon EKS and ECS



Fast, scalable, and secure

Performance/latency under high-load, policies with custom traffic and security plug-ins



AWS AI native integrations

Work with Amazon Bedrock as LLM and AI agent provider to optimize and API to AI API conversions

Developer API Platform / Automation/APIOps

Automate everything consistently with CLI, REST API, and YAML accessible runtime configuration CI/CD

Kong AI Gateway



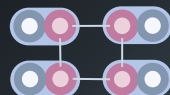
Kong Gateway



Kong Ingress Controller



Kong Mesh



AI use cases have come a long way in a short amount of time

Predict

Assist

Collaborate

Pioneer

2023

The year of POCs

2024

The year of Production

2025

The year of Business Value



What is generative AI?

Is this secure?



How do I prioritize my projects?



How can I leverage agents?

How can we go deeper *and* wider?

Do I need to become a prompt engineer?



How can I lower my costs?

How do I make this real?

How can we transform our entire business?



How do I choose a model?

Where do I get started?

What customization method should I use?



How can I get *all* my employees to embrace AI?



Enterprises are doubling down on Agentic AI

50%

Agentic AI Usage
Surged in 3 months

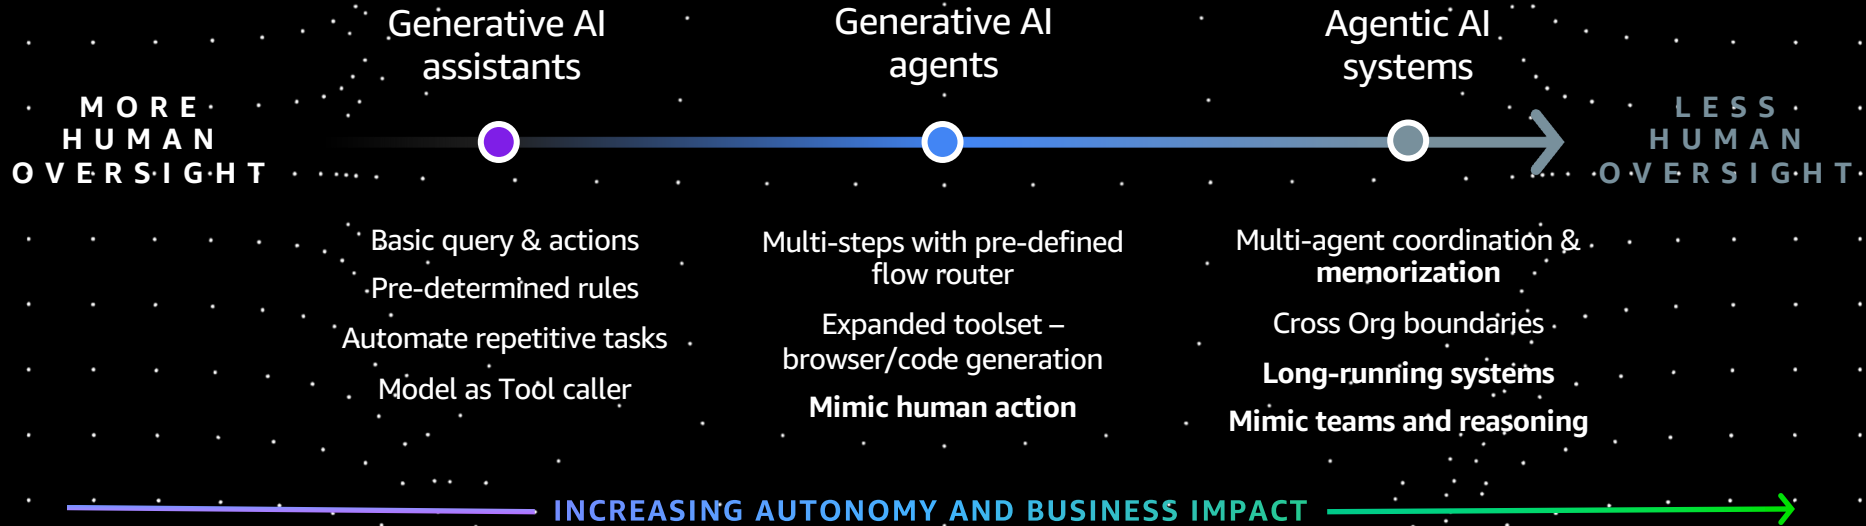
ACCORDING TO ADOBE, "Australia's agentic AI usage is accelerating fast", July 2025

217%

Higher Customer Satisfaction
with AI Agents

ACCORDING TO SALESFORCE, "Australian customers [...] more likely to prefer resolving service issues with an agent.[...]", October 2025

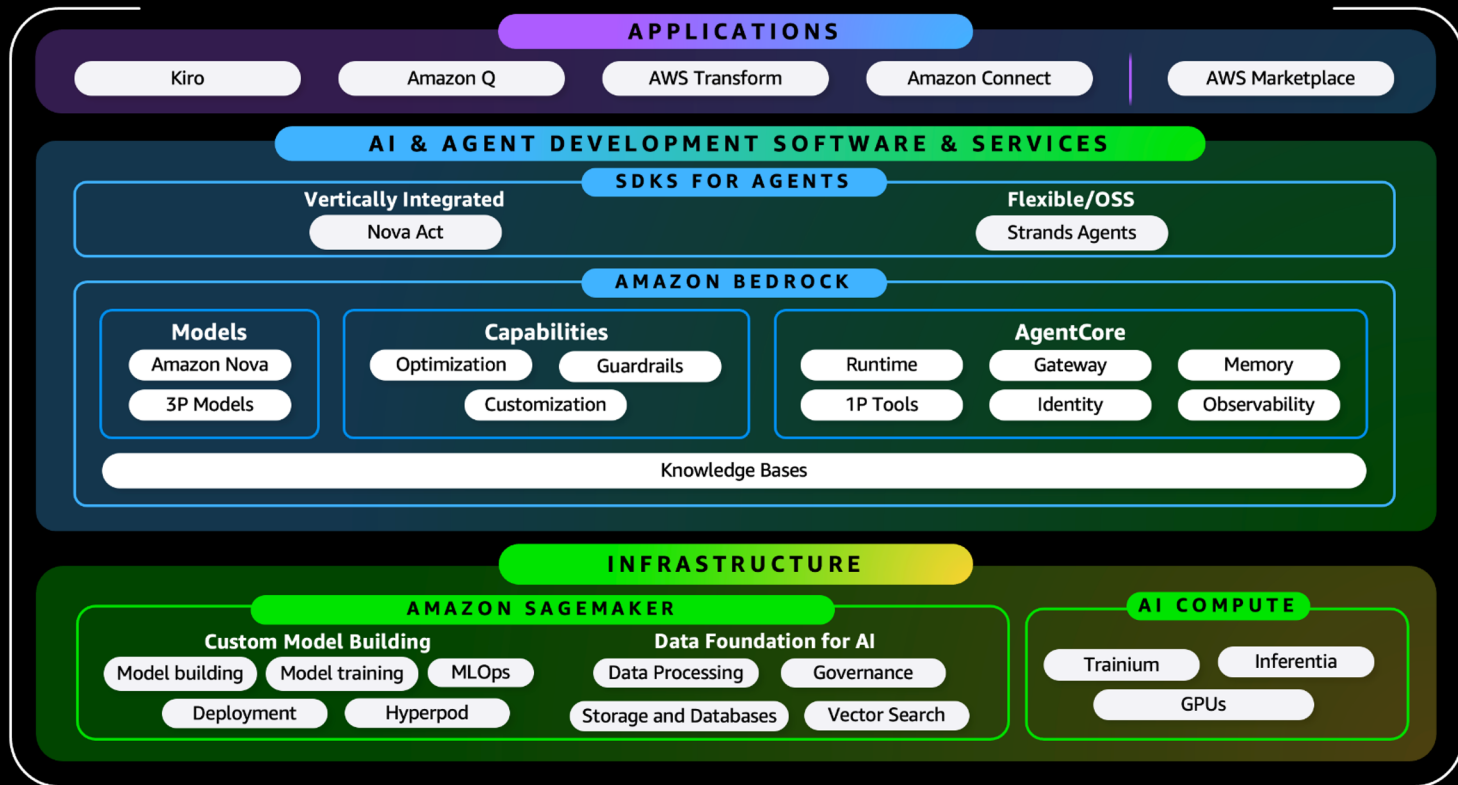
The evolution into Agentic AI



AWS Agentic AI portfolio

INTERFACES & PROTOCOLS
(MCP/A2A)

SECURITY & POLICIES



Amazon Bedrock

The easiest and fastest way to build and scale AI applications and agents

BROAD CHOICE OF
LEADING MODELS

CUSTOMIZE WITH
YOUR DATA

APPLY SAFETY &
RESPONSIBLE
AI CHECKS

OPTIMIZE FOR
COST, LATENCY &
ACCURACY


BUILD &
ORCHESTRATE
AGENTS

← AWS SECURITY, PRIVACY, AND RELIABILITY BUILT-IN →

At AWS,
everything starts with
security

SECURE BY DESIGN, AT EVERY LAYER, AT EVERY STAGE.

With AI, Security Risks and Threats remain (and emerge) at all layers



Responsible AI Dimensions

FAIRNESS

Considering impacts on different groups of stakeholders

EXPLAINABILITY

Understanding and evaluating system outputs

CONTROLLABILITY

Having mechanisms to monitor and steer AI system behavior

SAFETY

Preventing harmful system output and misuse

PRIVACY AND SECURITY

Appropriately obtaining, using, and protecting data and models

GOVERNANCE

Incorporating best practices into the AI supply chain, including providers and deployers

TRANSPARENCY

Enabling stakeholders to make informed choices about their engagement with an AI system

VERACITY AND ROBUSTNESS

Achieving correct system outputs, even with unexpected or adversarial inputs

Where Kong and Bedrock come together



Secure your API – Authentication / Authorization

OIDC with Amazon Cognito

Key Auth

AI Request & Response Transformer

Choose best Bedrock model, Customize with Agents/RAGs/MCPs

AI RAG Injector

AI Proxy Chat Route

AI Proxy: streaming, multi-LLM, files, batch

AI MCP Proxy

Optimize for cost, latency, and accuracy

AI Prompt Decorator, Template, Compressor

AI LLM as Judge

AI Semantic Cache

AI Rate Limiting Advanced

Safety and responsible AI checks

AI PII Sanitizer

AI Bedrock Guardrails

AI Prompt Guard

AWS Backend Services



Compute – containers,
VM workloads

Amazon Bedrock



Amazon Bedrock
LLM Models, Agents, RAGs

Bedrock Input
guardrails


Bedrock Output
guardrails



Amazon Bedrock
Guardrails

Implementing responsible AI policies

Guardrails - model probabilistic and/or deterministic and agnostic safeguards





Optimize for cost, latency, and accuracy



INFERENCE

Flexible and reduced-price
consumption options



MODELS

Latency-optimized models

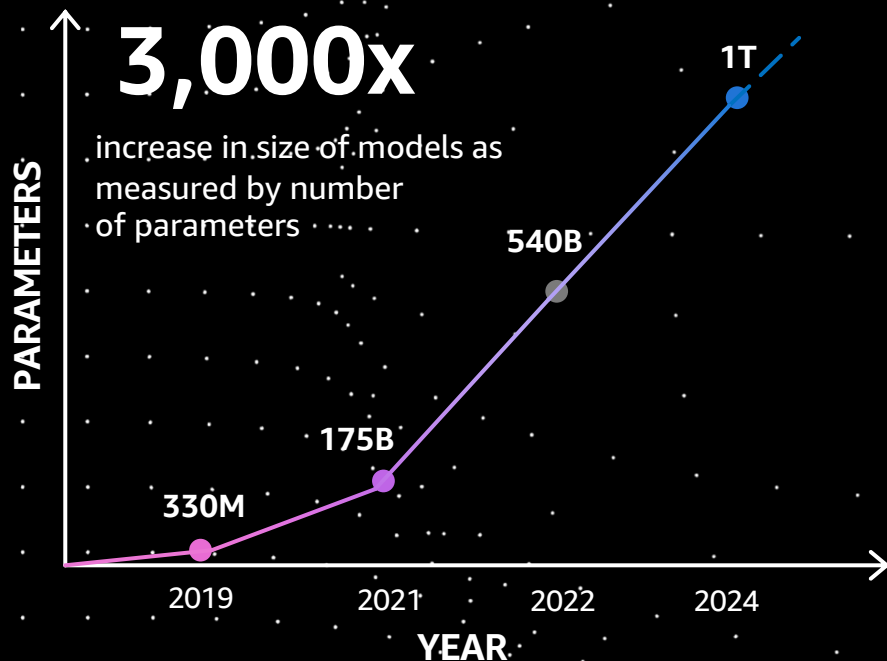


FEATURES

Model Distillation, Intelligent
Prompt Routing and Prompt
Caching

AI Models keep getting bigger

... A LOT BIGGER – LEADING TO INCREASED NEED FOR ACCELERATORS



The rise of large language models (LLMs) forces to **cost control carefully**

The rise of small language models (SLMs) and common optimization techniques increase the applicability of CPUs for GenAI inference

Optimization tactics in multiple places:

- Quantization
- RAG - Retrieval-Augmented Generation
- Prompt engineering
- Model compression - Pruning, Distillation

Amazon Bedrock Foundation Models (FMs)

BROADEST SELECTION OF FULLY MANAGED MODELS FROM LEADING AI COMPANIES

AI21 labs

Highly efficient processing
& grounded generation
for long context lengths

JAMBA

amazon

Frontier intelligence &
industry leading price
performance

NOVA

ANTHROPIC

Excels at complex
reasoning, code
generation, & instruction
following

CLAUDE

cohere

Powering efficient,
multilingual AI agents with
advanced search & retrieval

COMMAND
EMBED
RERANK

deepseek

Advanced reasoning
models that solve complex
problems step-by-step

DEEPSEEK-R1

Luma

High-quality video
generation with natural,
coherent motion &
ultra-realistic details

RAY2

Meta

Advanced image &
language reasoning

LLAMA

MISTRAL AI

Specialized expert models
for agentic reasoning and
multimodal tasks

MISTRAL
MIXTRAL
PIXTRAL

poolside (Coming soon)

Software engineering AI
for large enterprises

MALIBU
POINT

stability.ai

Professional-grade images
with creative control,
deployable at scale

STABLE DIFFUSION

TwelveLabs

CTRL + F for video data:
unlock the full potential of
enterprise video assets

MARENGO
PEGASUS

WRITER

Purpose-built models for
building & scaling AI
agents across the
enterprise

PALMYRA

Amazon Bedrock Marketplace

100+ PUBLICLY AVAILABLE AND PROPRIETARY EMERGING, POPULAR, AND SPECIALIZED MODELS
THROUGH A SERVERFUL OFFERING



IBM



NVIDIA



HUGGING FACE



MISTRAL AI



STABILITY AI



LG AI RESEARCH



SNOWFLAKE



UPSTAGE



PREFERRED
NETWORKS



CAMB.AI



WRITER



EVOLUTIONARY
SCALE



ARCEE AI



WIDN



GRETEL



DATABRICKS



KARAKURI



LIQUID



JOHN SNOW
LABS



CYBERAGENT



NCSoft

Amazon Bedrock inference consumption options

ON DEMAND

Pay-as-you-go,
no commitment

Pricing based on input and output token
count for LLMs

Great for prototyping, POCs, and small
workloads with more relaxed requirements
for throughput and latency

Requests per minute (RPM) and tokens per
minute (TPM) limits enforced

Includes support for **cross-region inference**

PROVISIONED THROUGHPUT

Provision sufficient throughput to
meet your application's
performance requirements

Reserve throughput (input/output tokens
per minute) at a **fixed cost**

Flexible commitment term of **1 month or
6 months**

Hourly PT (2MU), no commit, across
different models

Pay hourly rate, discounted for extended
commit

Great for **production workloads or
inference on custom models**

BATCH INFERENCE

Use Amazon Bedrock to
process prompts in batch to
get responses for model
evaluation, experimentation,
and offline processing.

Amazon Bedrock offers select foundation
models (FMs) from leading AI providers like
Anthropic, Meta, Mistral AI, and Amazon for
batch inference at 50% of on-demand
inference pricing

Efficiently **run model inference on large
volumes of data**

Avoids throttling when running
large jobs

Amazon Bedrock Cost Optimization

Amazon Bedrock Model Distillation

Utilize **smaller, more cost-effective models**

Maximize distilled model performance with proprietary data synthesis

Distilled models are **up to 500% faster and up to 75% less expensive** than original models, with less than 2% accuracy loss for use cases like RAG

Amazon Bedrock Intelligent Prompt Routing

Automatically route prompts to different foundation models to optimize response quality and lower costs

Provides a single endpoint to efficiently route prompts

Uses advanced prompt matching techniques to meet cost and latency thresholds

Reduce costs by up to 30%

Kong and Amazon Bedrock support for prompt caching

Semantically cache repetitive context in prompts across multiple API calls

Enhance accuracy through longer, detailed prompts

Compress prompts by reducing length while preserving meaning

Reduce costs by up to 90% and latency by up to 85% for supported models

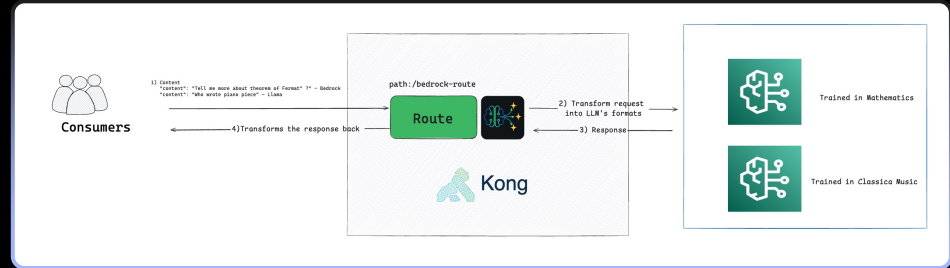
Kong AI Built-in Cost Optimization Features

AI Rate limiting

Use the token data returned by the LLM provider to calculate the costs of queries.

Allow to analyze and **assign costs to incoming queries, then rate limit the consumer's cost** for a given time window and provider.

Intelligent Routing



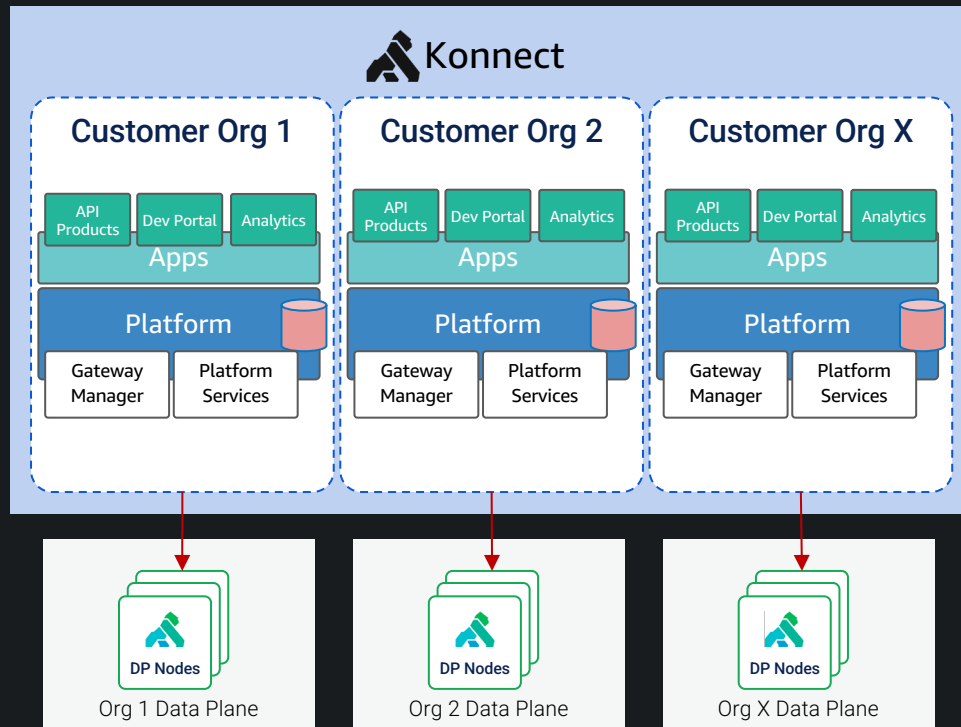
Prompt Compression

Ratio-based or target token compression — for example, reduce a message to 80% of the original length or compress to 150 tokens.

Kong Konnect - Gateway Manager Run Through

Konnect SaaS Multi-tenancy

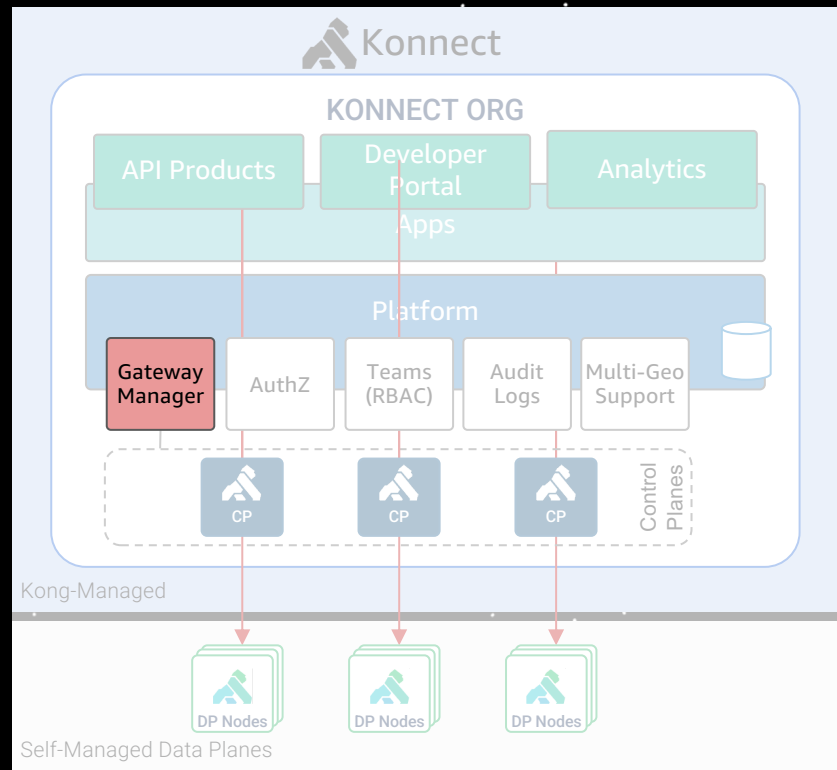
- Konnect is multi-tenanted - multiple orgs share the platform.
- Control planes are provided as a service to each organization.
- Data planes are managed by the tenant outside of Konnect.
- Tenants are isolated via common SaaS strategies (AWS Best Practice) and certified as SOC-2 Type 2 and CSA Star-Level 1 compliant.
- Konnect guarantees 99.9% uptime SLA.
- Konnect downtime won't disrupt API traffic.



What is the Konnect Gateway Manager?

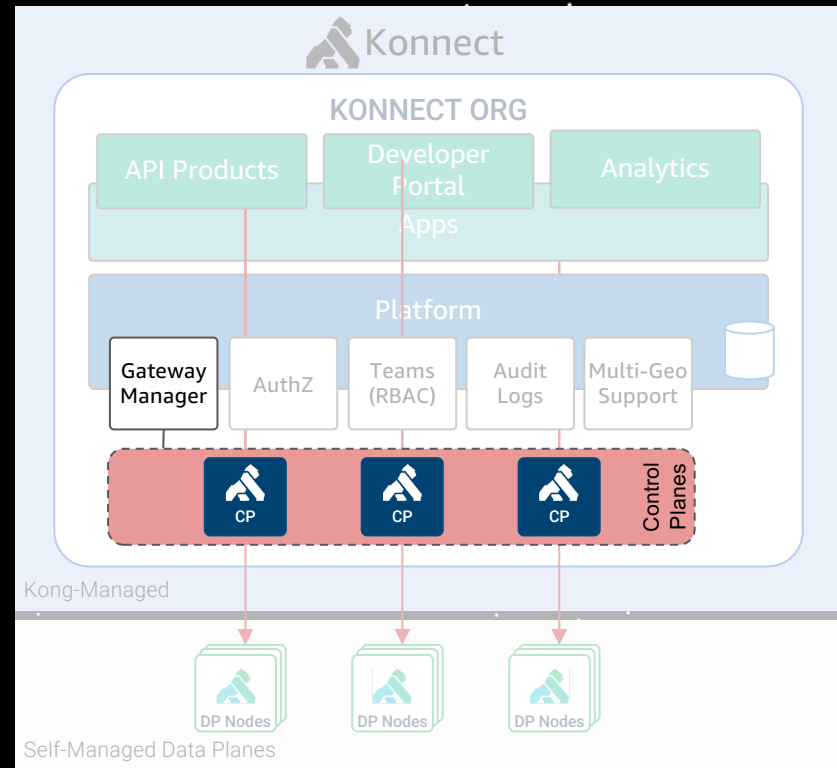
Gateway Manager allows you to:

- Catalogue, connect, & monitor the status of all control planes and their associated data plane nodes.
- Add or delete user defined control planes (Default control plane can be renamed but not deleted).
- Add, delete, and manage data plane nodes and their configuration.
- Securely isolate API gateway instances, e.g. dev/staging/prod environments, business units, etc.

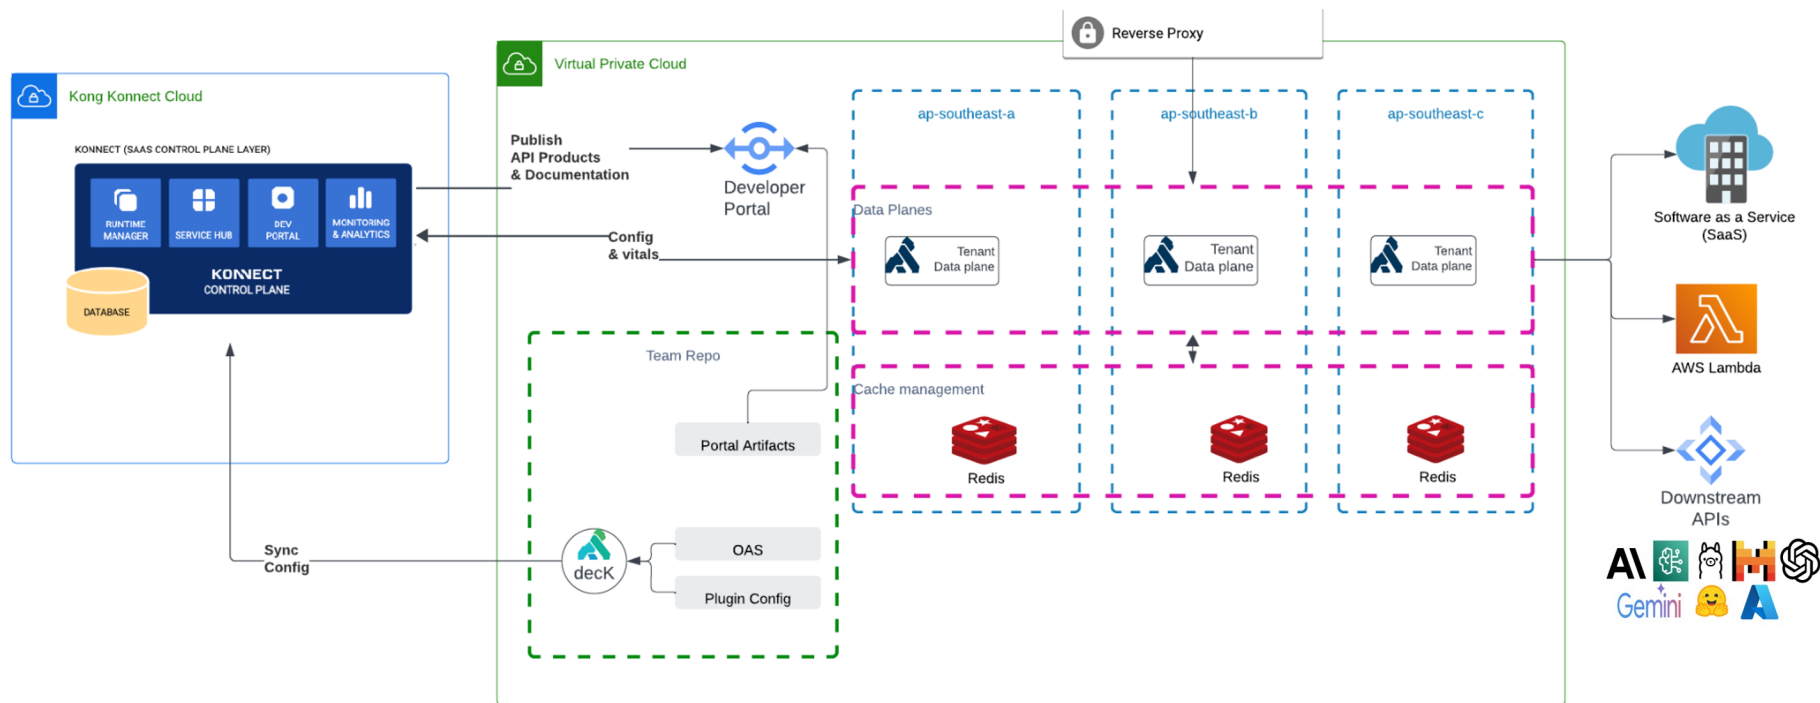


What is a Konnect Control Plane?

- A Control Plane is hosted in the cloud by Kong and offers centralized management, configuration, and monitoring of APIs.
- Each control plane stores the configuration for its data plane nodes, but does not have access or visibility into the application data flowing through them.
- Each Konnect Organization can have multiple control planes, each with its own configuration and telemetry data.
- Categories of control planes:
 - Kong Gateway (Hybrid)
 - Kong Ingress Controller
 - Composite Control Plane



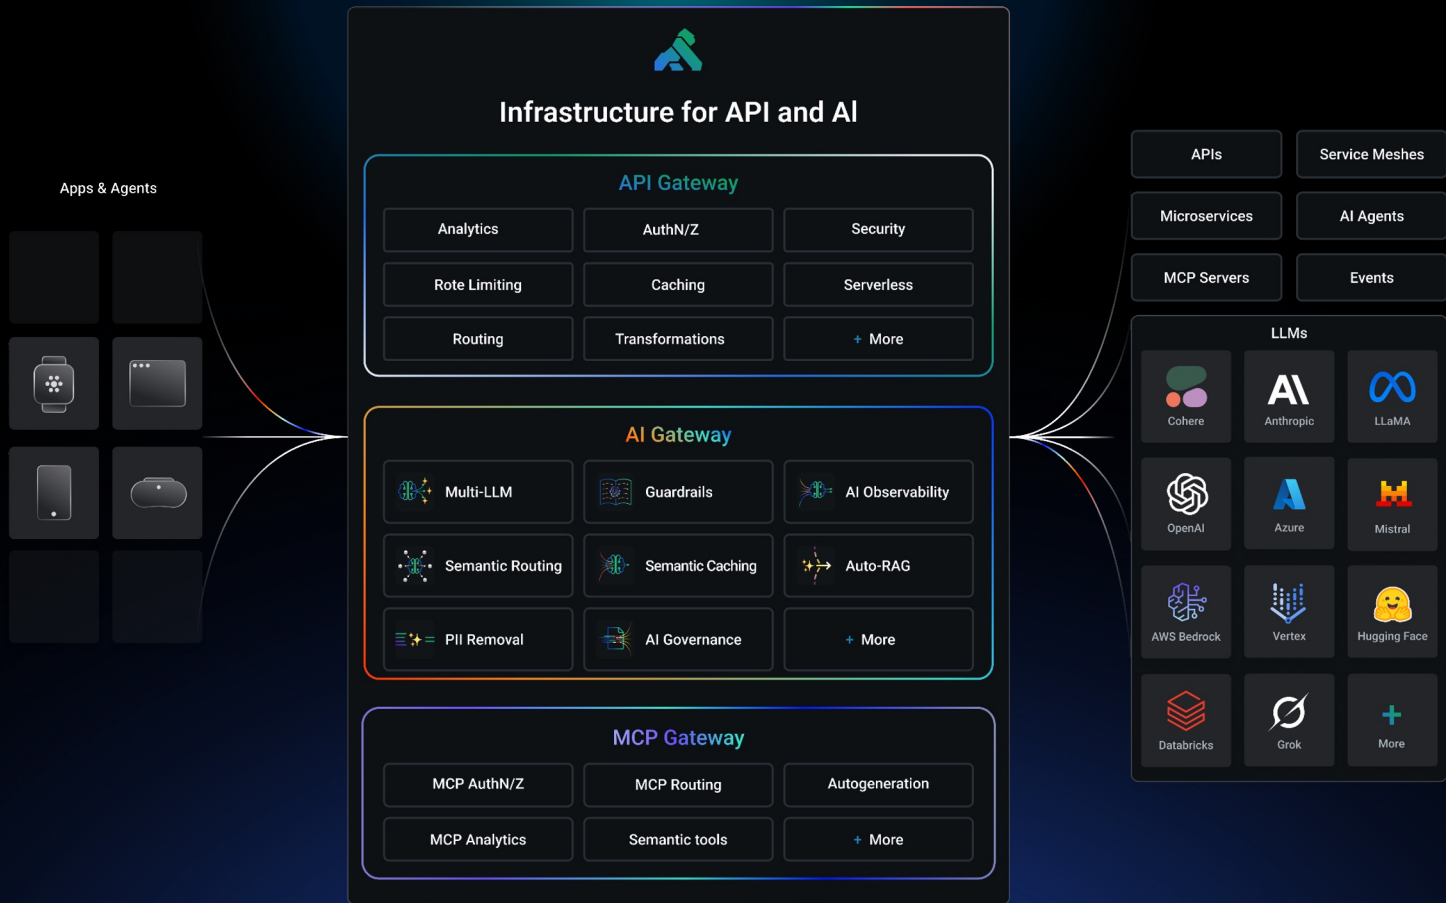
Deploy with AWS



There is no AI without APIs

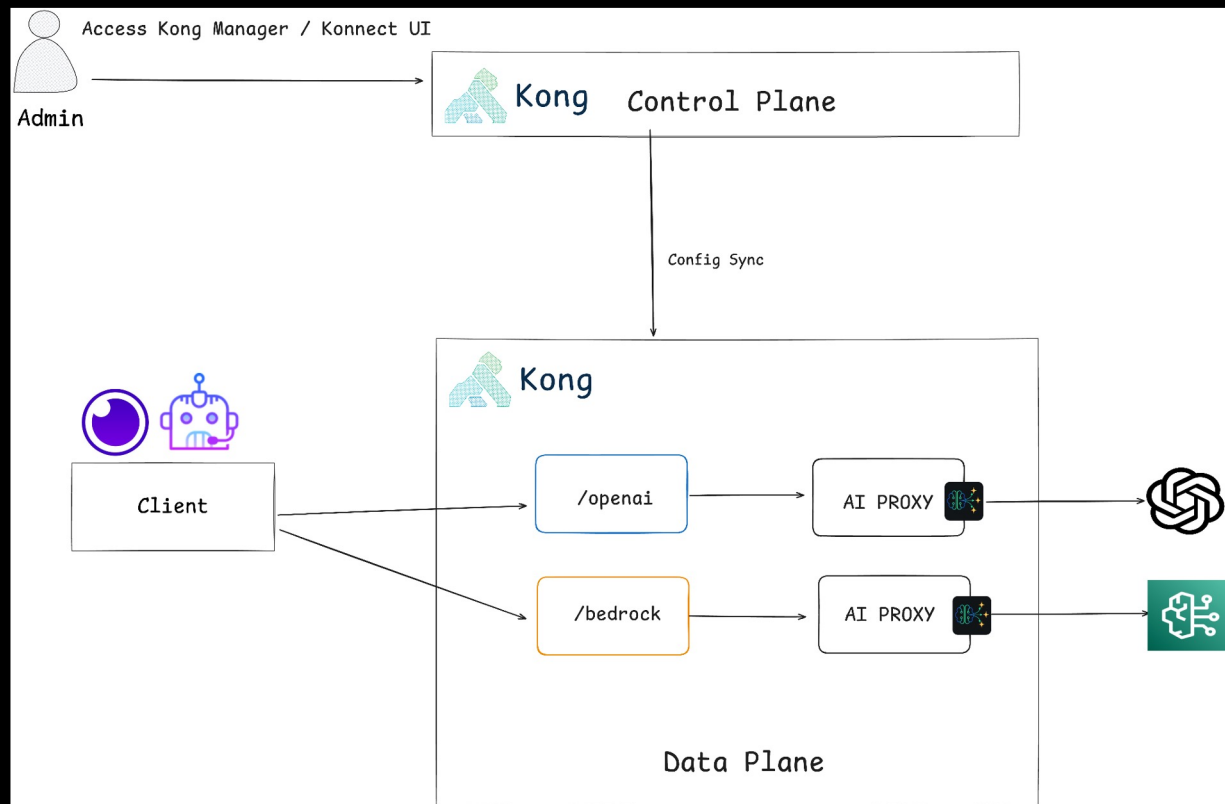


Kong Must Power The AI Agentic Infrastructure



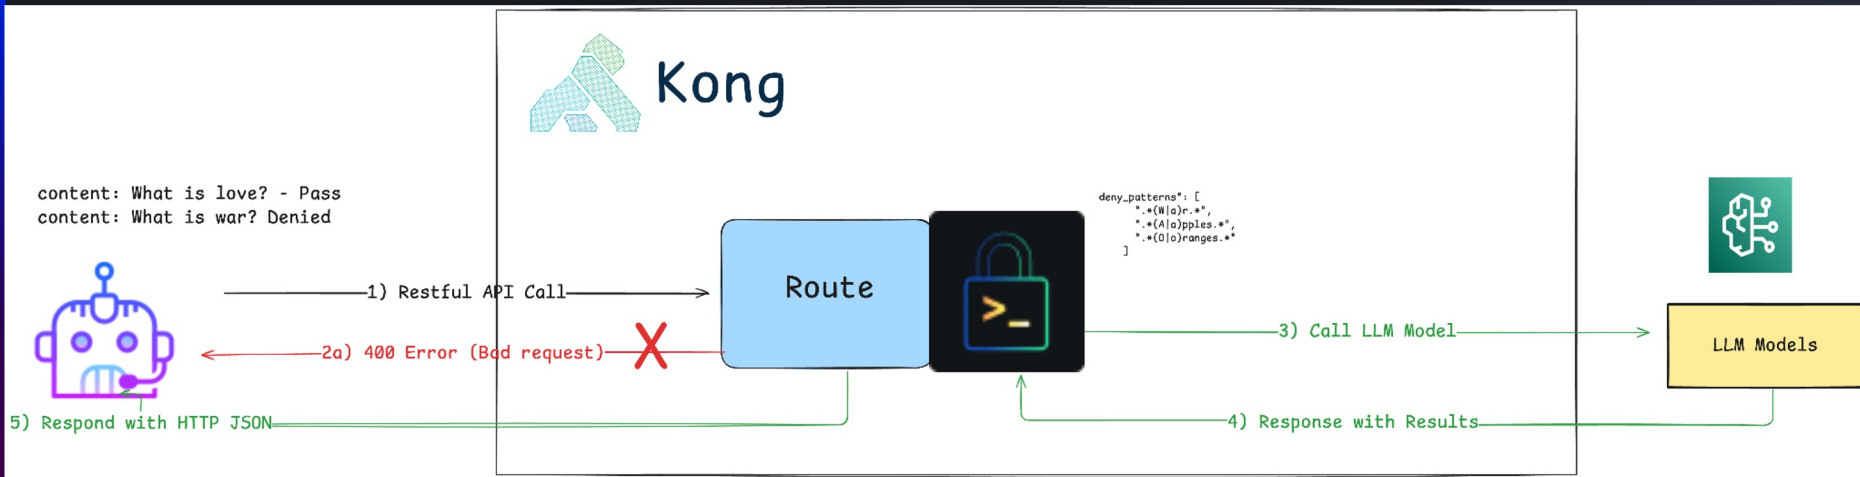
Use Cases

Proxy to Bedrock

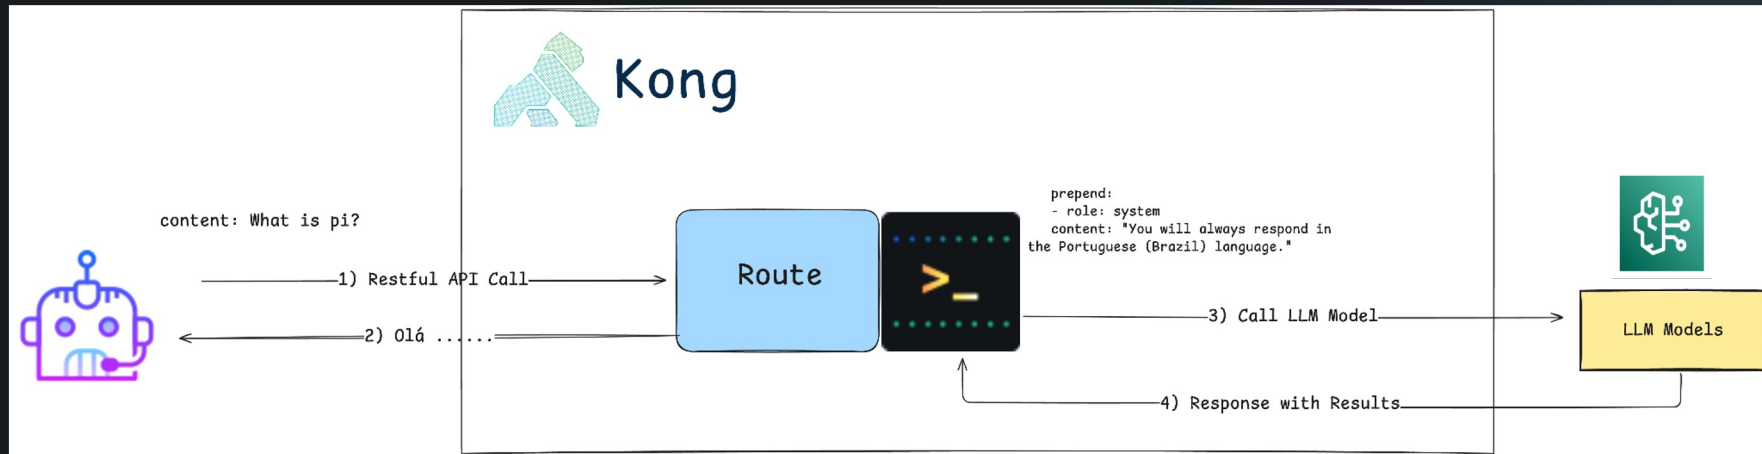


Prompt Management

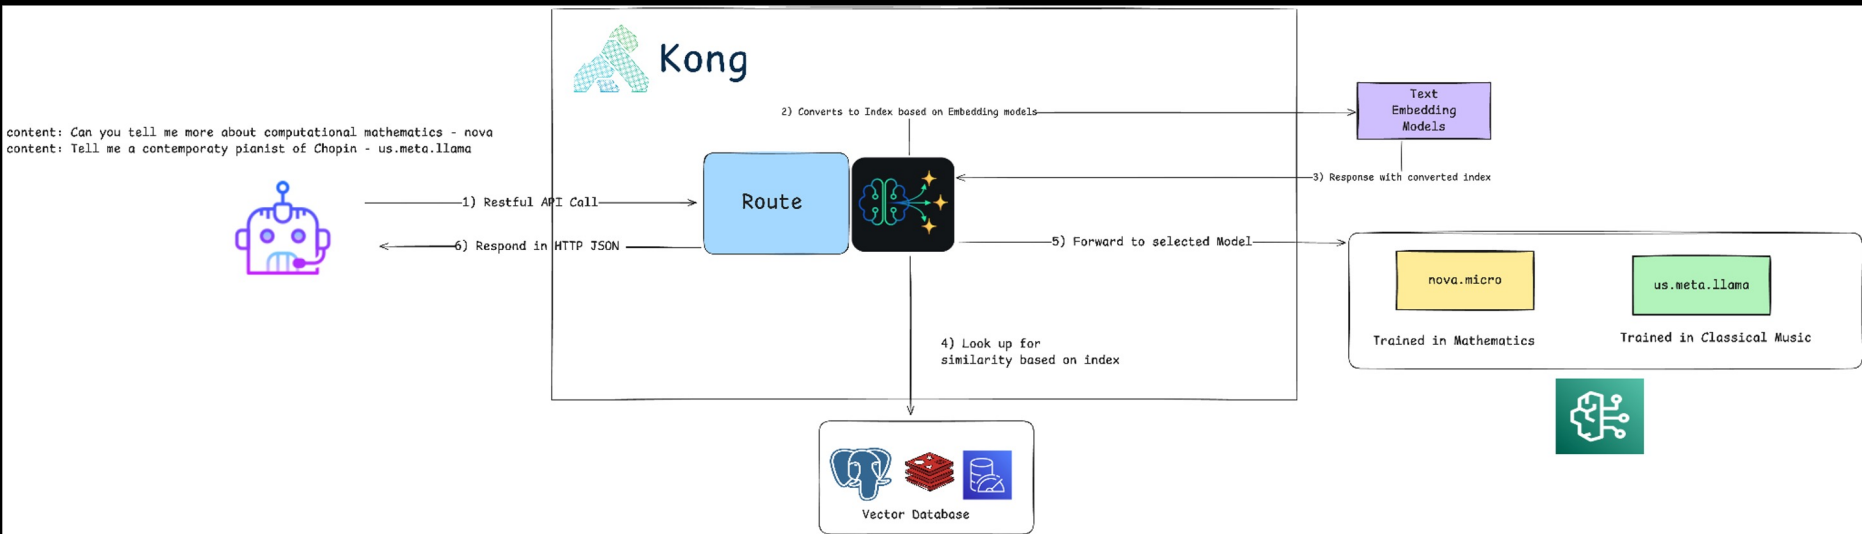
Prompt Guard



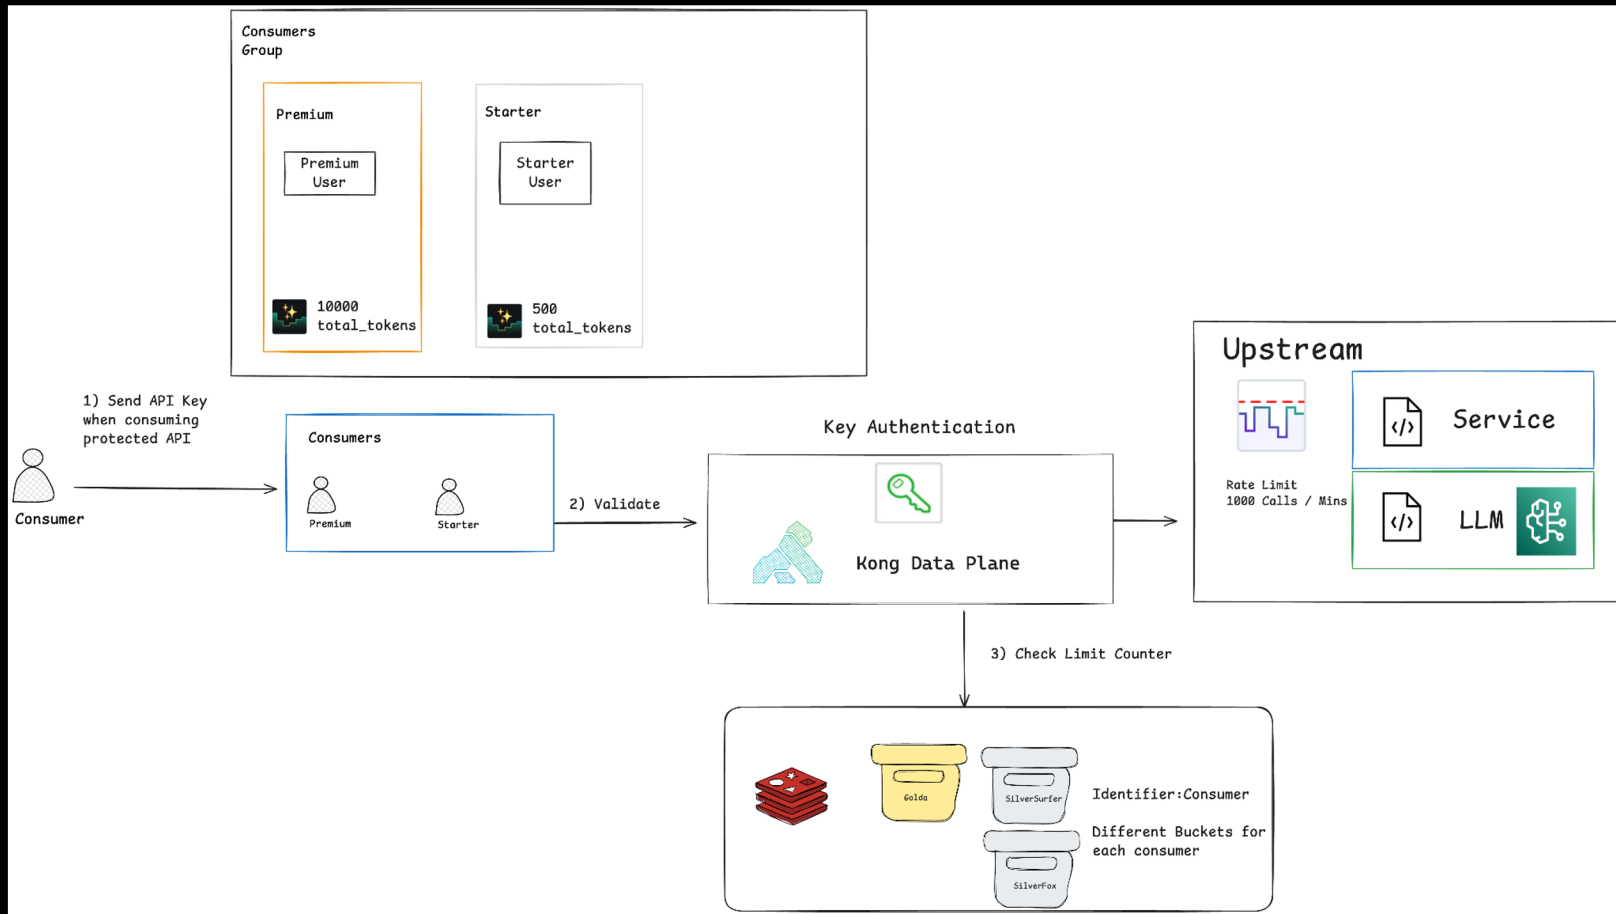
Prompt Decorator



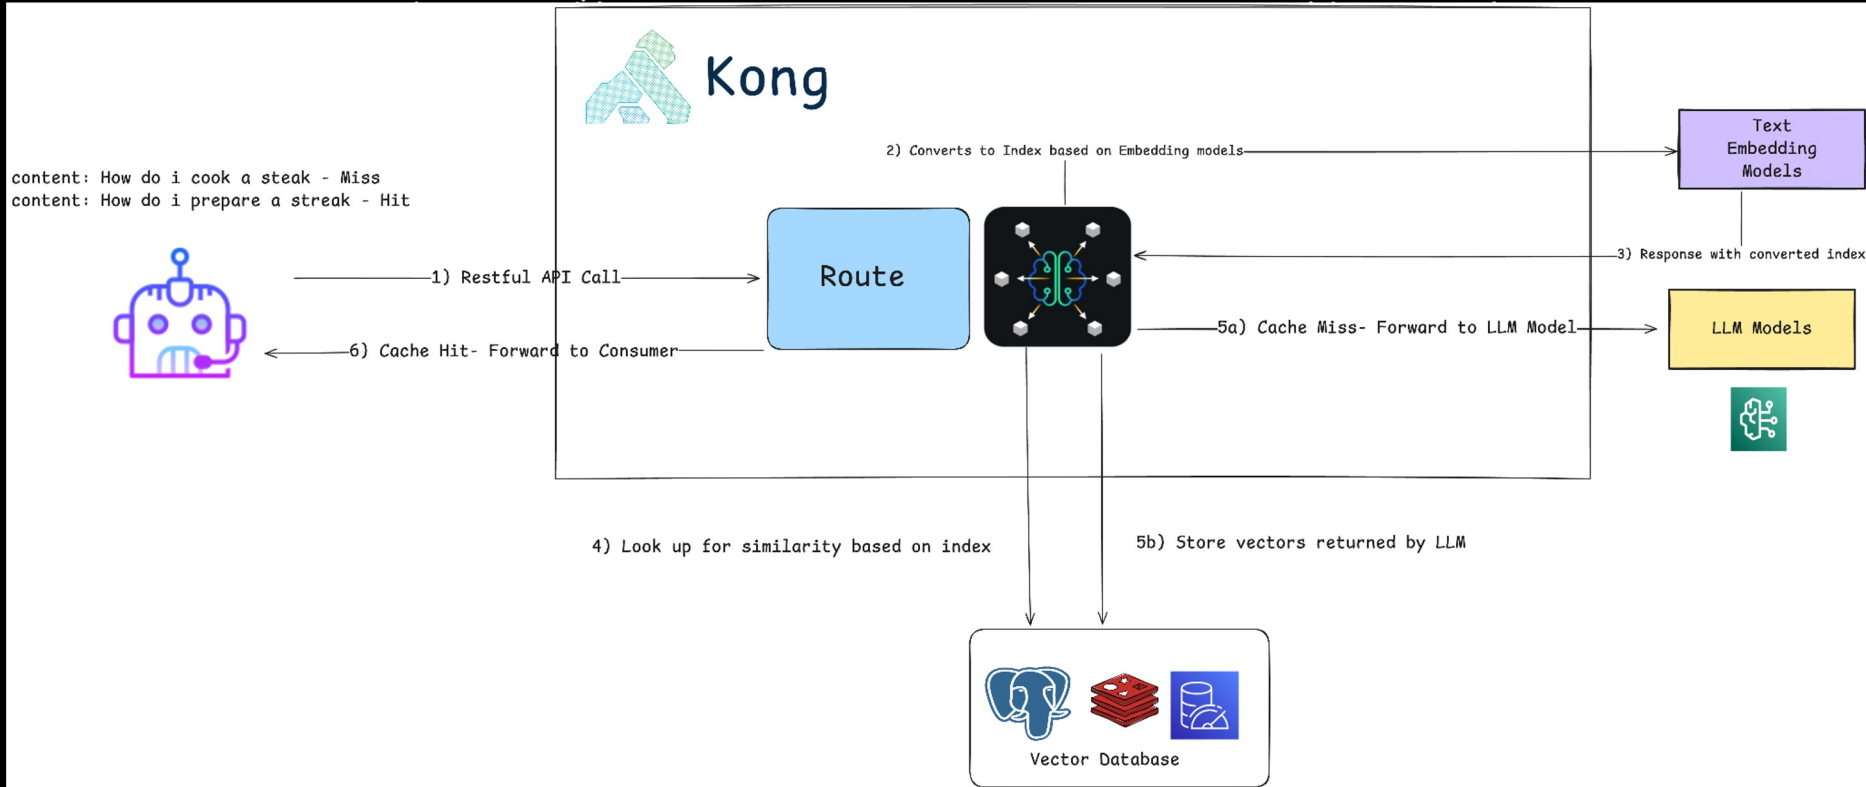
Semantic Routing



AI Budgeting

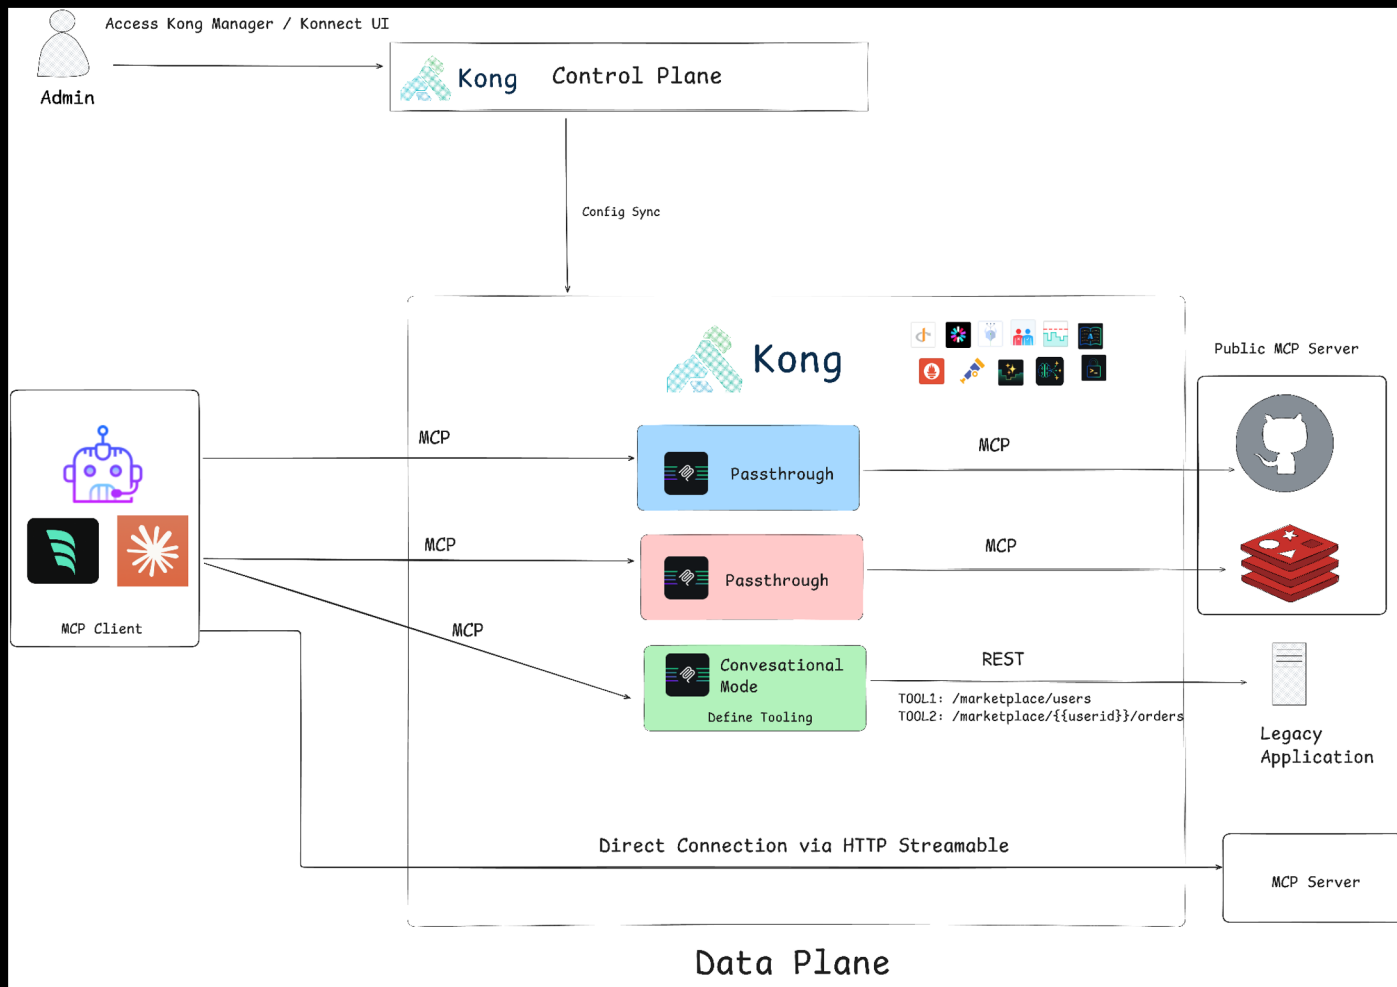


AI Semantic Caching



Indexes previous prompts and responses as embeddings. On each request, it searches for semantically similar inputs and serves cached responses when possible to reduce redundant LLM calls.

AI MCP Proxy



Kong **Konnect** Cloud

API Portal

API & AI Analytics

Gateway & Mesh Manager

Service Hub Service Registry

Kong **Insomnia**

Build & Consume

API Design

API Testing

API Client

AI Apps

Kong **Gateway** Ingress & AI

Edge - Internal - AI

APIs

API Gateway


Kube Ingress

AI Gateway

AI LLMs

Kong **Mesh**

Within Apps




Self managed data planes
Hybrid Gateway, AI & Mesh

 Cloud &
On-Premise

SaaS dedicated data planes
API Gateway, AI Gateway

 20+ Regions
Autoscaling





Amazon Bedrock AgentCore

Deploy and operate highly capable agents securely, at scale using any framework and model

TIME TO VALUE



Build powerful AI agents without the infrastructure and operational headaches

FLEXIBLE



Create agents with any framework or model

TRUSTED



Deploy, secure, scalable and reliable agents your organization can trust

Build custom agents with Bedrock AgentCore



Open source

Protocols:
MCP, A2A

Frameworks:
CrewAI, LangGraph, Kong
Kong Volcano,
AWS Strands Agents



Amazon Bedrock



AgentCore



Runtime



Identity



Gateway



Code
interpreter



Memory



Browser
tool



Observability

Model Capabilities



Model access



Cost and performance
optimization




Customization




Guardrails

Knowledge Bases

What an Agentic AI system looks like with Kong and Bedrock AgentCore



Start Now!





Connect and reach out if we can help!


1.

Kong and Bedrock AWS Workshop



2.

AWS Generative AI Innovation Center



3.

Kong in AWS Marketplace





KONG AND AWS PARTNERCAST

Robin Cher

Solutions Architect
At Kong

robin.cher@konghq.com

Adrian Phang

Cloud Ecosystem Lead
At Kong

adrian.phang@konghq.com

William Garcia

Solutions Architect
At Amazon Web Services

willgrod@amazon.com

Cassandra Se To

Partner Development Manager
At Amazon Web Services

caseto@amazon.com



AWS PARTNERCAST

Thank you!

Please join us again for another PartnerCast session

<https://aws.amazon.com/partners/training/partnercast/>